

A Framework for Feature Selection in Clustering


Daniela M. Witten and Robert Tibshirani

Stanford University



Wentao Wu

Problem

- Clustering: Given $X_{n \times p}$, identify K clusters.
- Sparse Clustering: Underlying clusters only differ on $q < p$ features.
 - improve accuracy and interpretation
 - cheaper prediction



Motivating Example



$$X_1 \sim N \left[\begin{pmatrix} -0.5 \\ 0 \end{pmatrix}, \begin{pmatrix} 1 & 0.78 \\ 0.78 & 1 \end{pmatrix} \right]$$

$$X_2 \sim N \left[\begin{pmatrix} 0.5 \\ 0 \end{pmatrix}, \begin{pmatrix} 1 & 0.78 \\ 0.78 & 1 \end{pmatrix} \right]$$

```
1 #K-means
2 cl <- kmeans(x,2)
3 plot(x,col=cl$cluster, main='Standard K-
  means')
4 #Sparse K-means
5 km.perm <- KMeansSparseCluster.permute(x,K
  =2,wbounds=seq(3,7,len=15),nperms=5)
6 km.out <- KMeansSparseCluster(x,K=2,wbounds=
  km.perm$bestw)
7 plot(x,col=km.out[[1]]$Cs, main='Sparse K-
  means')
```

Sparse Clustering

$$\begin{aligned} & \max_{\Theta \in D} \quad \sum_{j=1}^p w_j f_j(X_j, \Theta) \\ \text{subject to} \quad & \|w\|^2 \leq 1, \quad \|w\|_1 \leq s, \quad w_j \geq 0 \end{aligned}$$

- w_j is a weight corresponding to feature j , also indicates contribution.
- $w_1 = w_2 = \dots = w_p$, reduces to the traditional clustering.
- L_1 penalty results in sparsity.
- L_2 penalty avoids trivial solution, (at most one element of w is nonzero).

How to optimize

- Alternating Iterative Algorithm
- Holding w fixed, optimize with respect to Θ .
 - Standard clustering procedure to a weighted version of the data.
- Holding Θ fixed, optimize with respect to w .
 -

$$\begin{aligned} & \max_w w^T a \\ \text{subject to} \quad & \|w\|^2 \leq 1, \|w\|_1 \leq s, w_j \geq 0 \end{aligned}$$

How to optimize

$$\begin{aligned} & \max_w \quad w^T a \\ & \text{subject to} \quad \|w\|^2 \leq 1, \quad \|w\|_1 \leq s, \quad w_j \geq 0 \end{aligned}$$

Introducing Lagrangian multipliers λ_1 for $\|w\|^2 - 1 \leq 0$, λ_2 for $\|w\|_1 - s \leq 0$ and μ_j for $w_j \geq 0$. We obtain the KKT condition:

$$\begin{aligned} w^T w - 1 &\leq 0 & \sum w_j - s &\leq 0 \\ \mu_j w_j &\leq 0 & w_j &\geq 0 \\ \lambda_1 (w^T w - 1) &= 0 & \lambda_2 (\sum w_j - s) &= 0 \\ \mu_j w_j &= 0 & -a + \lambda_1 w_j + \lambda_2 - \mu_j &= 0 \end{aligned}$$

How to optimize

Based on the last equation, $\mu_j = -a + \lambda_1 w_j + \lambda_2$. Thus, $(-a + \lambda_2 + \lambda_1 w_j)w_j = 0$.

- Case 1: $\sum w_j - s < 0$. Thus, $\lambda_2 = 0$. If $a < 0$, then $w_j = 0$. Otherwise, $w_j = \frac{a}{\lambda_1}$, where λ_1 make sure that $w^T w = 1$.
- Case 2: $\sum w_j - s = 0$. Thus, $\lambda_2 \geq 0$. If $a < 0$, then $w_j = 0$. Otherwise, $w_j = \frac{a - \lambda_2}{\lambda_1}$, where λ_1 ensures that $w^T w = 1$ and λ_2 ensures $\sum w_j - s = 0$.

In summary,

$$w = \frac{S(a_+, \Delta)}{\|S(a_+, \Delta)\|^2}$$

- $a = f_j(X_j, \Theta)$.
- $S(x, c) = \text{sign}(x)(|x| - c)_+$.
- $\Delta = 0$ if that results in $\|w\|_1 \leq s$, otherwise, $\Delta > 0$ is chosen to yield $\|w\|_1 = s$.

How to optimize

```
1 FindDelta <- function(a,l1bound){
2   #Find optimal delta which statisfy w constrains
3
4   #if already statisfies, return 0
5   if(l2n(a)==0 || sum(abs(a/l2n(a)))<=l1bound) return(0)
6   #binary search lo is 0, hi is maximum absolute element of a
7   delta.lo <- 0
8   delta.hi <- max(abs(a))-1e-5
9   iter <- 1
10  while(iter<=15 && (delta.hi-delta.lo)>(1e-4)){
11    #cal S(a,delta)
12    su <- soft(a,(delta.lo+delta.hi)/2)
13    if(sum(abs(su/l2n(su)))<l1bound){
14      #if stasifies the l1 bound, try a smaller delta
15      delta.hi <- (delta.lo+delta.hi)/2
16    } else {
17      #if not, try last delta
18      delta.lo <- (delta.lo+delta.hi)/2
19    }
20    iter <- iter+1
21  }
22  return((delta.lo+delta.hi)/2)
23 }
```


Sparse K-Means

$$\begin{aligned} \max_{C_1, C_2, \dots, C_K, w} \quad & \sum_{j=1}^p w_j \left(\frac{1}{n} \sum_{i=1}^n \sum_{i'=1}^n d_{i,i',j} - \sum_{k=1}^K \frac{1}{n_k} \sum_{i,i' \in C_k} d_{i,i',j} \right) \\ \text{subject to} \quad & \|w\|^2 \leq 1, \quad \|w\|_1 \leq s, \quad w_j \geq 0 \end{aligned}$$

- w is the weight for each feature.
- K is the number of clusters.
- n_k is the number of observations in cluster k .
- C_k contains the indices of the observations in cluster k .
- $d_{i,i',j}$ is the dissimilarity measure between observation i and i' along feature j .

Optimize Sparse K-Means Clustering

- ① Initialize w as $w_i = \frac{1}{\sqrt{p}}$, where $i = 1, \dots, p$.
- ② Iterative until convergence $\frac{\sum_{j=1}^p |w_j^r - w_j^{r-1}|}{\sum_{j=1}^p |w_j^{r-1}|} < 10^{-4}$.
 - ① Holding w fixed, optimize with respect to C_1, \dots, C_K . Applying the standard K-means algorithm to the $n \times n$ dissimilarity matrix with (i, i') element $\sum_j w_j d_{i, i', j}$.
 - ② Holding C_1, \dots, C_K fixed, optimize with respect to w by applying

$$w = \frac{S(a_+, \Delta)}{\|S(a_+, \Delta)\|^2}$$

where $a_j = \frac{1}{n} \sum_{i=1}^n \sum_{i'=1}^n d_{i, i', j} - \sum_{k=1}^K \frac{1}{n_k} \sum_{i, i' \in C_k} d_{i, i', j}$

- ③ The clusters are given by C_1, \dots, C_K and the feature weights corresponding to this clustering are given by w_1, \dots, w_p .

Optimization

```
1 UpdateC <- function(x, K, w, C=NULL){
2   #num of row
3   n <- nrow(x)
4   #remove zero-weight features
5   x <- x[,w!=0];
6   xr <- sweep(x, 2, sqrt(w[w!=0]), "*");
7   if(!is.null(C)){
8     #num of column
9     p <- ncol(x)
10    #attach cluster information
11    xr <- data.frame(data=xr, clu=C)
12    #calculate cluster mean
13    mu <- aggregate(xr,by=list(xr$clu),FUN=mean)
14    #cal new dist including the cluster mean
15    xr <- as.matrix(xr[,1:p])
16    mu <- mu[,1:p+1]
17    mu <- as.matrix(mu)
18    distmat <- as.matrix(dist(rbind(xr, mu)))[1:n, (n+1):(n+K)]
19    #get the cluster
20    c_n <- apply(distmat, 1, which.min)
21    #do the clustering
22    if(length(unique(c_n))==K){
23      kc <- kmeans(xr,centers=mu)
24    }else{
25      kc <- kmeans(xr,centers=K, nstart=10)
26    }
27  }else{
28    kc <- kmeans(xr, centers=K, nstart=10)
29  }
30 }
```

Optimization

```
1 UpdateW <- function(x, C, l1bound){  
2   #calculate WCSS  
3   wcss.feature <- CalWCSS(x, C)$wcss.feature  
4   #calculate TSS  
5   tss.feature <- CalWCSS(x, rep(1, nrow(x)))$wcss.feature  
6   #find delta  
7   delta <- FindDelta(-wcss.feature+tss.feature, l1bound)  
8   #soft operation  
9   wu.unscaled <- soft(-wcss.feature+tss.feature,delta)  
10  return(wu.unscaled/l2n(wu.unscaled))  
11 }
```

Tuning Parameter





- s the L_1 bound on w .
- Use gap statistics.
- Gap statistic measures the strength of the clustering obtained on the real data relative to the clustering obtained on null data that does not contain subgroups.

Tuning Parameter

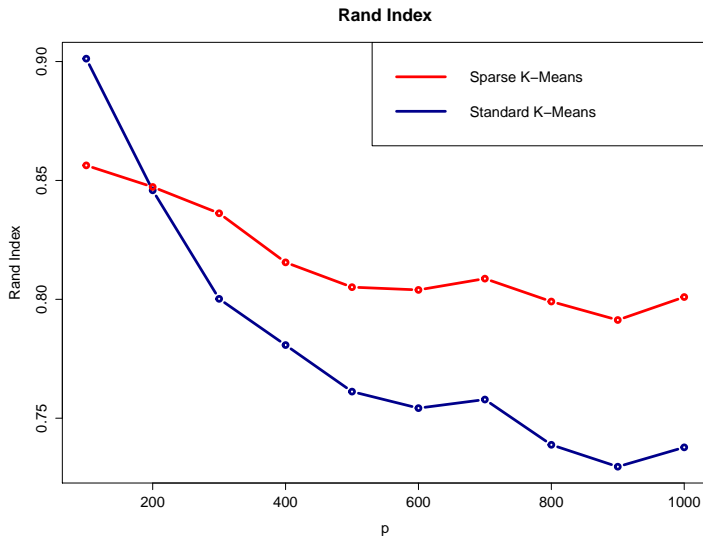
- ① Obtain permuted datasets X_1, \dots, X_B by independently permuting the observations within each feature.
- ② For each candidate tuning parameter value s :
 - ① Compute $O(s)$ the objective obtained by performing sparse K-means with tuning parameter s on the data X .
 - ② For $b = 1, \dots, B$, compute $O_b(s)$.
 - ③ Calculate $gap(s) = \log(O(s)) - \frac{1}{B} \sum_{b=1}^B \log(O_b(s))$.
- ③ Choose s^* as the smallest value for which $gap(s^*)$ is within a standard deviation of the largest value of $gap(s)$.

- Three Clusters: C_1, C_2, C_3 .
- Each cluster contains 20 observations.
- $p = 100, q = 50$: 100 features, the underlying clusters depend on the first 50 features.
- $X_{ij} \sim N(\mu_{ij}, 1)$.
- If $i \in C_k$ and $j \leq q$, $\mu_{ij} = \mu_{C_k}$, where $\mu_{C_1} = 0, \mu_{C_2} = 0.6, \mu_{C_3} = 1.2$.
- If $j > q$, $\mu_{ij} = 0$ regardless of i .
- Metric: Rand index

Simulation-Different S




Simulation-Different p



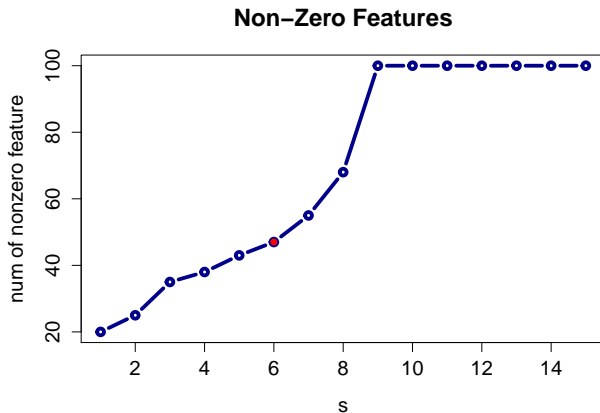
Simulation-Tuning

```
1 k.perm <- SparseKMeansTuning(x,K=3,wbounds=seq(3,10,len=15), nperm=10)
```



Simulation-Tuning

```
1 k.perm <- SparseKMeansTuning(x,K=3,wbounds=seq(3,10,len=15), nperm=10)
```



Sparse Hierarchical

$$\begin{aligned} & \max_U \quad \{ \sum_j \sum_{i,i'} w_j d_{i,i',j} U_{i,i'} \} \\ \text{subject to} \quad & \sum_{i,i'} U_{i,i'}^2 \leq 1, \quad \|w\|^2 \leq 1, \quad \|w\|_1 \leq s, \quad w_j \geq 0 \end{aligned}$$

- $U_{i,i'} \propto \sum_j w_j d_{i,i',j}$.
- Performing hierarchical clustering on U results in the sparse hierarchical clustering.

How to optimize

- ➊ Initialize w as $w_1 = \dots = w_p = \frac{1}{\sqrt{p}}$.
- ➋ Iterate until convergence:
 - ➊ Update $u = \frac{Dw}{\|Dw\|_2}$.
 - ➋ Update $w = \frac{S(a_+, \Delta)}{\|S(a_+, \Delta)\|_2}$ where $a = D^T u$ and $\Delta = 0$ if this results in $\|w\|_1 \leq s$; otherwise, $\Delta > 0$ is chosen such that $\|w\|_1 = s$.
- ➌ Rewrite u as a $n \times n$ matrix U .
- ➍ Perform hierarchical clustering on the $n \times n$ dissimilarity matrix U .

The End